OBJECT DETECTION · SET PREDICTION · REAL TIME

NMS-Free Real-Time Detection: DETR → RT-DETR

Move duplicate suppression from post-processing into learning—and redesign the compute bottleneck

Central claim — DETR learns duplicate suppression through one-to-one Hungarian matching. RT-DETR retains this principle while redesigning the encoder, query learning, and decoder for real-time inference.

42.0	53.1	108	34.8
DETR R50 AP	RT-DETR R50 AP	FPS · T4 · TensorRT FP16	RT-DETR AP_S · small-object gap

Problem — Anchors · duplicate boxes · NMS cost › **DETR** — Set prediction · optimal one-to-one matching › **RT-DETR** — Efficient encoder · better queries · tunable depth

01 · STARTING POINT What changed in DETR?

Conventional detectors generate many proposals or anchors, then use heuristics and NMS to remove duplicates. DETR predicts the image's object set in one pass.

Image backbone features › Transformer global context › 100 learned queries parallel decoder › Object set classes + boxes

Matching before the loss

The Hungarian algorithm finds the globally optimal one-to-one assignment between prediction slots and ground-truth objects. Each object is assigned to only one slot, training the model to produce a duplicate-free set without NMS.

$$L = L_{class} + \lambda_1 L_1 + \lambda_G L_{GIoU}$$

L1 directly penalizes coordinate error; GIoU supplies overlap and relative-geometry information. Under the same 15% relative error, practical check [B] measured a large/small L1 ratio of **9.44×**, versus **1.000×** for GIoU.

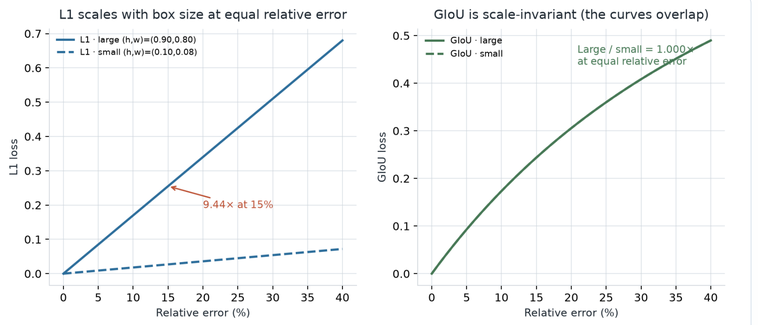


Fig. 1. L1 and GIoU response to box scale — practical check [B], synthetic boxes.

02 · THE NMS DILEMMA Accuracy–latency controls

In RT-DETR Table 1's YOLOv8 results, raising the confidence threshold removes more candidates: NMS gets faster, but AP drops. Raising the IoU threshold suppresses fewer boxes: AP rises, but processing takes longer.

Control	AP	NMS latency
conf 0.001 → 0.05	52.9 → 51.2	2.36 → 1.06 ms
IoU 0.5 → 0.8	52.1 → 52.8	2.24 → 2.46 ms

NMS-free inference removes this threshold trade-off from post-processing and relies on one-to-one matching during training. Since YOLO speed reports commonly exclude NMS, the paper establishes an end-to-end benchmark that includes the NMS plugin.

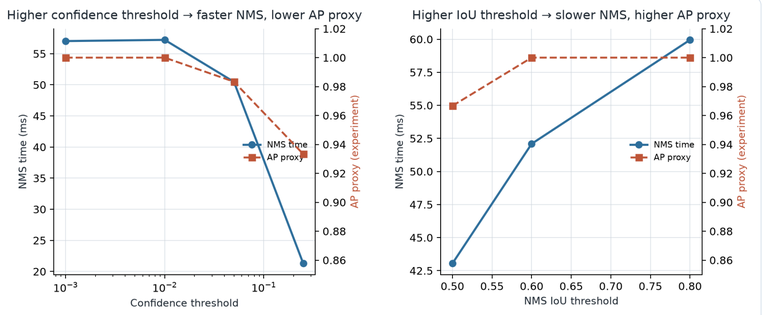


Fig. 2. Practical check [A]: CPU, synthetic data. Raising conf 0.001→0.25 reduced time 57.03→21.29 ms and AP proxy 1.000→0.933. Absolute values are not comparable to the paper; only the direction is.

Boundary: The paper does not evaluate Video Object Detection, cross-frame identity association, or frame-level IoU. NMS-free detection is not evidence that temporal tracking is solved.

03 · PERFORMANCE CONTEXT DETR by object scale

DETR has similar overall AP to Faster R-CNN-FPN+, yet its scale-wise profile differs: stronger on large objects and weaker on small objects.

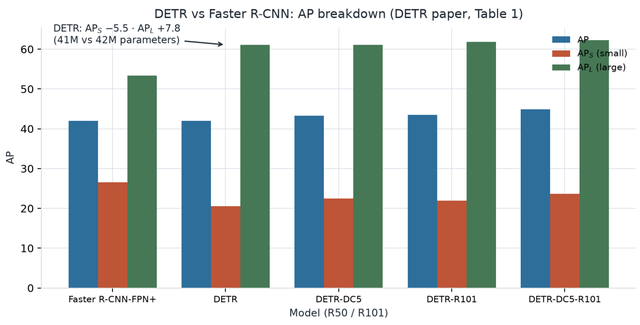


Fig. 3. DETR vs Faster R-CNN — scale breakdown from paper Table 1. AP_S −5.5; AP_L +7.8.

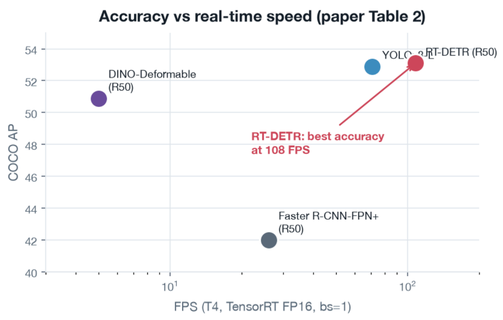


Fig. 4. Accuracy vs real-time speed — paper Table 2. RT-DETR is the only model in the top-right: highest AP and 108 FPS on T4 (TensorRT FP16, bs=1). NMS time is included for the YOLO baselines.

04 · BOTTLENECK → DESIGN Why is AIFI applied only to S5?

RT-DETR §3.2 cites a Deformable-DETR encoder accounting for **49% of GFLOPs** but only **11% of AP contribution**. This is a second-hand statistic from prior work—not a measurement of RT-DETR itself.

Table 3 variant	AP	Latency
A · remove encoder	43.0	7.2 ms
B · single-scale Transformer	44.9	11.1 ms
C · multi-scale Transformer	45.6	13.3 ms
D · decouple interaction/fusion	46.4	12.2 ms
D_{S5} · attention on S5 only	46.8	7.9 ms
E · AIFI + CCFF	**47.9**	**9.3 ms**

D: separating intra-scale interaction from cross-scale fusion is faster and more accurate than C. **D→D_{S5}:** limiting attention to S5 cuts latency by about 35% and adds 0.4 AP.

AIFI applies self-attention to low-resolution, semantically rich S5 features; CNN-based CCFF fuses scales. The authors argue that lower-level features lack semantic concepts and risk *"duplication and confusion"* with high-level interactions.

$$\text{Attention pairs} \propto M^2$$
$$(8400 / 400)^2 = 441\times$$

This is a theoretical token-pair comparison, not a 441× end-to-end latency claim. In practical check [B], 8× more tokens (400→3200) increased measured attention time by 42.3×.

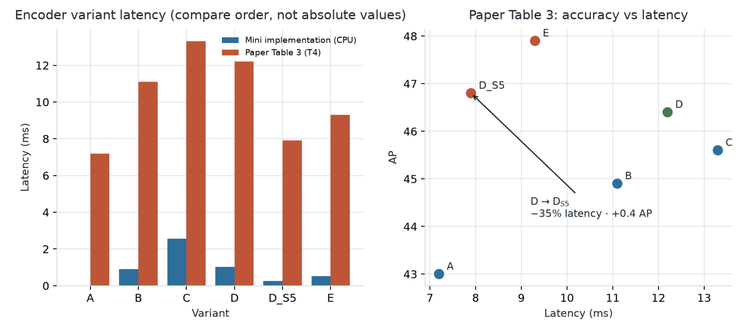


Fig. 4. Practical check [C] reproduces the paper's latency order: A < D_{S5} < E < B < D < C. Absolute runtimes are not compared across hardware.

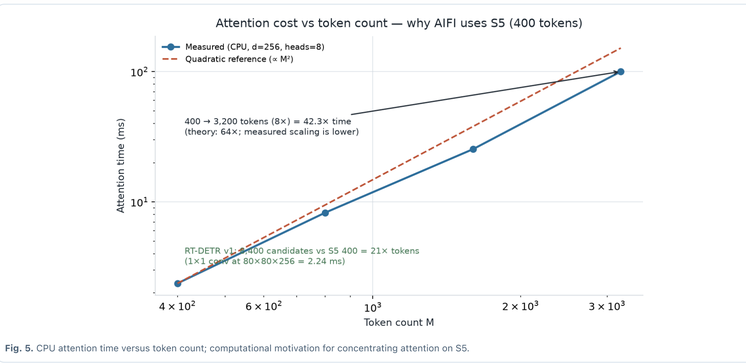


Fig. 5. CPU attention time versus token count; computational motivation for concentrating attention on S5.

05 · QUERY QUALITY What does uncertainty change?

$$U(X) = \|P_{location}(X) - C_{class}(X)\|$$

Equation (2) measures disagreement between localization and classification predictions. Equation (3) puts U in the training loss so encoder features used as queries learn better-aligned class and localization quality.

47.9 → 48.7	0.35 → 0.82%	0.30 → 0.67%
AP · +0.8	Prop cls	Prop both

In Fig. 6, purple points concentrate in the upper-right, where both classification and IoU scores are high. The reported 138% and 120% are **relative count differences** under score thresholds—not percentage-point gains.

Ranking ≠ loss: Public code still ranks top-K by classification score. The key change is not inference-time ranking; it is training the encoder so that features selected by that ranking are higher quality.

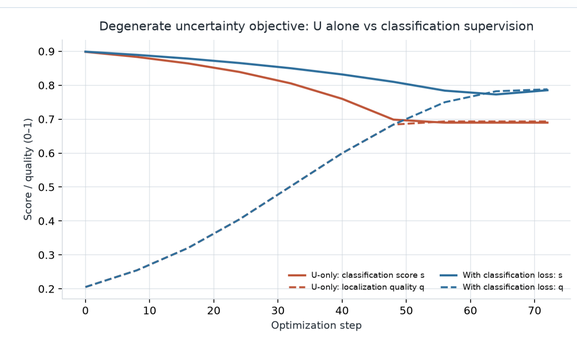


Fig. 6. In practical check [D], optimizing U alone led to the degenerate solution s 0.90→0.690 and q 0.20→0.693. Adding classification supervision avoids collapsing quality into mere score agreement.

06 · INFERENCE CONTROL Tune speed with decoder depth

Auxiliary prediction losses train intermediate decoder layers to predict, so inference can stop before the final layer.

Det^6 output	AP	Latency
6 layers	53.1	9.3 ms
5 layers	53.0	8.8 ms
1 layer	49.1	7.0 ms

Each layer adds roughly 0.4–0.5 ms, while AP gains shrink from +2.2 (layer 1→2) to +0.1 (5→6). For deployment, `eval_idx` chooses the decoder output; `max_det` separately caps returned detections.

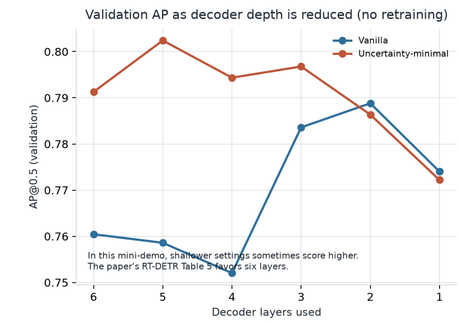


Fig. 7. Decoder depth provides a no-retraining inference-speed control.

The Day 07 mini-demo sometimes showed slightly higher AP at 5 layers than at 6. Its synthetic data, few objects, and short training differ from the paper's COCO experiment: structural reproduction is not reproduction of the AP curve.

07 · RESULTS Faster and more accurate

Model	AP	AP_S	AP_L	FPS
DETR R50	42.0	20.5	61.1	28
RT-DETR R50	53.1	34.8	70.0	108

Conditional context: RT-DETR uses 640×640 and 72 epochs; DETR uses 800×1333 and 500 epochs. Different input sizes and schedules mean these are not controlled absolute speed/accuracy comparisons.

08 · OPEN LIMITS What remains unresolved?

RT-DETR-R50 AP_S is 34.8, 0.5 below YOLOv8-L at 35.3; R101 is 0.9 below its comparison model. The 49/11 statistic is second-hand, FPS depends on hardware/runtime, and the paper evaluates 2D COCO—not 3D depth estimation or multi-view detection.

09 · STUDY PROCESS Claim → evidence → boundary

We checked original tables and equations, then used matching, NMS, attention, and query-loss experiments to test directional claims. These CPU/synthetic checks are not reproductions of COCO/GPU paper results.

Takeaway: Move duplicate suppression into learning, compute into the architecture, and initialization into quality.

References and materials

Carion et al., "End-to-End Object Detection with Transformers," ECCV 2020, arXiv:2005.12872 · Lv et al., "DETRs Beat YOLOs on Real-time Object Detection," CVPR 2024, arXiv:2304.08069.

Experiment charts were regenerated from the Day 07 DETR and Day 08 RT-DETR metrics; all chart labels and captions are in English. CPU/synthetic mini-demo measurements are distinguished from paper results.